

Map Reduce

Design Patterns

Caveat

- MapReduce is a **framework** not a tool
 - You have to fit your solution into the framework of map and reduce
 - It might be challenging in some situations
- Need to take the algorithm and break it into **filter/aggregate steps**
 - Filter becomes part of the map function
 - Aggregate becomes part of the reduce function
- Sometimes we may need **multiple** Map Reduce **stages**
- Map Reduce is **not a solution to every problem**, not even every problem that profitably can use many compute nodes operating in parallel!
- It makes sense only when:
 - files are **very large** and are **rarely updated**
 - We need to **iterate** over all the files to **generate** some interesting property of the data in those files

Design Patterns

- Intermediate data reduction
- Matrix generation and multiplication
- Selection and filtering
- Joining
- Graph algorithms

Intermediate Data

- Written locally
 - Transferred from mappers to reducers over network
- Issue
 - Performance bottleneck
- Solution
 - Reduce data
 - Use combiners
 - Use In-Mapper Combining

In-Mapper Combining (I)

```
1: class MAPPER
2:     method MAP(docid  $a$ , doc  $d$ )
3:         for all term  $t \in \text{doc } d$  do
4:             EMIT(term  $t$ , count 1)

1: class REDUCER
2:     method REDUCE(term  $t$ , counts [ $c_1, c_2, \dots$ ])
3:          $sum \leftarrow 0$ 
4:         for all count  $c \in \text{counts } [c_1, c_2, \dots]$  do
5:              $sum \leftarrow sum + c$ 
6:         EMIT(term  $t$ , count  $sum$ )
```

In-Mapper Combining (II)

```
1: class MAPPER
2:   method MAP(docid  $a$ , doc  $d$ )
3:      $H \leftarrow$  new ASSOCIATIVEARRAY
4:     for all term  $t \in$  doc  $d$  do
5:        $H\{t\} \leftarrow H\{t\} + 1$ 
6:     for all term  $t \in H$  do
7:       EMIT(term  $t$ , count  $H\{t\}$ )
```

In-Mapper Combining (III)

```
1: class MAPPER
2:   method INITIALIZE
3:      $H \leftarrow \text{new ASSOCIATIVEARRAY}$ 
4:   method MAP(docid  $a$ , doc  $d$ )
5:     for all term  $t \in \text{doc } d$  do
6:        $H\{t\} \leftarrow H\{t\} + 1$ 
7:   method CLOSE
8:     for all term  $t \in H$  do
9:       EMIT(term  $t$ , count  $H\{t\}$ )
```

In-Mapper Combining

- **Advantages:**

- **Complete** local aggregation **control** (how and when)
- **Guaranteed** to execute
- Direct **efficiency control** on intermediate data creation
- Avoid **unnecessary** objects creation and destruction (before combiners)

- **Disadvantages:**

- **Breaks** the functional programming background (state)
- Potential ordering-dependent **bugs**
- **Memory** scalability **bottleneck** (solved by memory **foot-printing** and flushing)

Matrix Generation

- Common problem:
 - Given an input of size N , generate an output of size $N \times N$
- Example: word co-occurrence matrix
 - Given a document collection, emit the bigram frequencies
- Two solutions
 - **Pairs**: generating $O(N^2)$ data in $O(1)$ space
 - **Stripes**: generation $O(N)$ data in $O(N)$ space

Pairs

```
1: class MAPPER
2:   method MAP(docid  $a$ , doc  $d$ )
3:     for all term  $w \in \text{doc } d$  do
4:       for all term  $u \in \text{NEIGHBORS}(w)$  do
5:         EMIT(pair ( $w, u$ ), count 1)                                ▷ Emit count for each co-occurrence

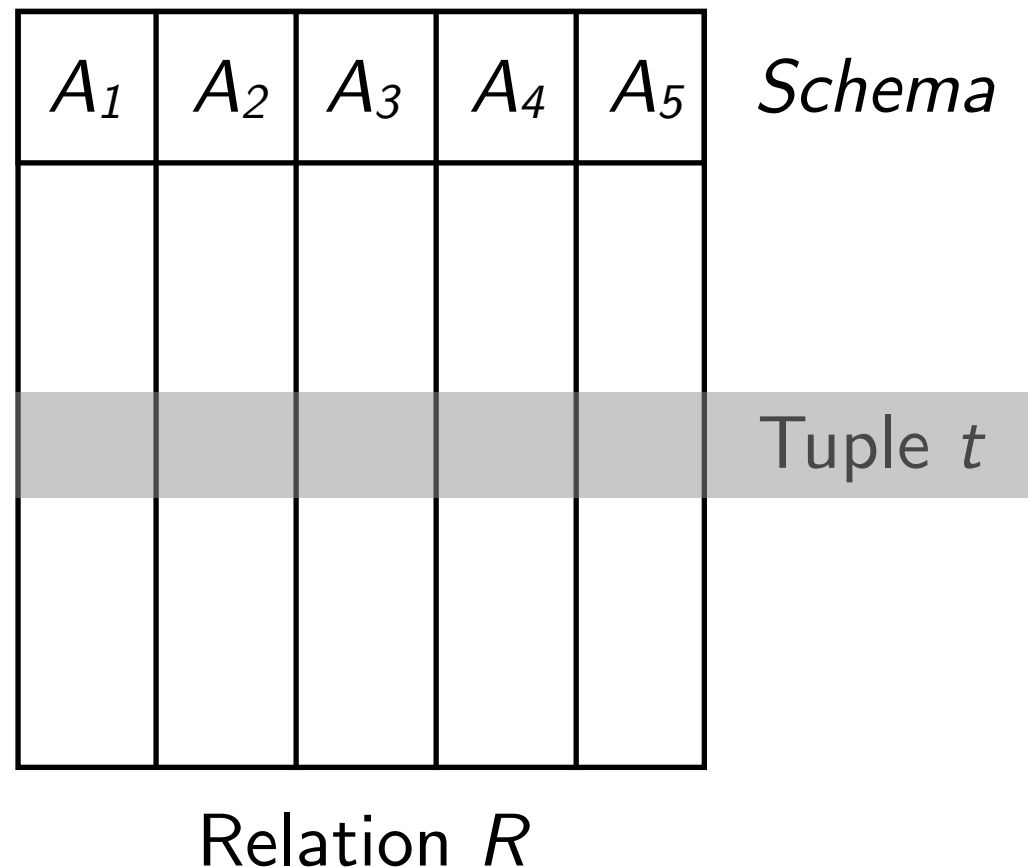
1: class REDUCER
2:   method REDUCE(pair  $p$ , counts [ $c_1, c_2, \dots$ ])
3:      $s \leftarrow 0$ 
4:     for all count  $c \in \text{counts } [c_1, c_2, \dots]$  do
5:        $s \leftarrow s + c$                                             ▷ Sum co-occurrence counts
6:     EMIT(pair  $p$ , count  $s$ )
```

Stripes

```
1: class MAPPER
2:   method MAP(docid  $a$ , doc  $d$ )
3:     for all term  $w \in \text{doc } d$  do
4:        $H \leftarrow \text{new ASSOCIATIVEARRAY}$ 
5:       for all term  $u \in \text{NEIGHBORS}(w)$  do
6:          $H\{u\} \leftarrow H\{u\} + 1$  ▷ Tally words co-occurring with  $w$ 
7:       EMIT(Term  $w$ , Stripe  $H$ )

1: class REDUCER
2:   method REDUCE(term  $w$ , stripes [ $H_1, H_2, H_3, \dots$ ])
3:      $H_f \leftarrow \text{new ASSOCIATIVEARRAY}$ 
4:     for all stripe  $H \in \text{stripes } [H_1, H_2, H_3, \dots]$  do
5:       SUM( $H_f, H$ ) ▷ Element-wise sum
6:     EMIT(term  $w$ , stripe  $H_f$ )
```

Relational Algebra Operators



- **SELECTION:** Select from relation R tuples satisfying condition $c(t)$
- **PROJECTION:** For each tuple in relation R , select only certain attributes A_i
- **UNION, INTERSECTION, DIFFERENCE:** Set operations on two relations with same schema
- **NATURAL JOIN**
- **GROUPING and AGGREGATION**

Selection and projection

- **Map:** each tuple t in R , if condition $c(t)$ is satisfied, is outputted as a (t, t) pair
 - **Reduce:** for each (t, t) pair in input, output (t, \perp)
-
- **Map:** each tuple t in R , create a new tuple t' containing only the projected attributes and output a (t', t') pair
 - **Reduce:** for each (t', t') pair in input, output (t', \perp)

Selection and projection

- **Map:** for each tuple t in R , if condition $c(t)$ is satisfied, output a (t, t) pair
 - **Reduce:** for each (t, t) pair in input, output (t, \perp)
-
- **Map:** for each tuple t in R , create a new tuple t' containing only the projected attributes and output a (t', t') pair
 - **Reduce:** for each $(t', [t', t', t', t'])$ pair in input, output (t', \perp)

Union, intersection and difference

- **Map:** for each tuple t in R , output a (t, t) pair
- **Reduce:** for each input key t , there will be 1 or 2 values equal to t . Coalesce them in a single output (t, \perp)
- **Map:** for each tuple t in R , output a (t, t) pair
- **Reduce:** for each input key t , there will be 1 or 2 values equal to t . If there are 2 value, coalesce them in a single output (t, \perp) otherwise do nothing
- **Map:** for each tuple t in R , output $(t, \textcolor{red}{R})$ and for each tuple t in S , output $(t, \textcolor{green}{S})$
- **Reduce:** for each input key t , there will be 1 or 2 values. If there is 1 value equal to $(t, \textcolor{red}{R})$ output (t, \perp) , otherwise do nothing

Natural Join

For simplicity, assume we have two relations $R(A,B)$ and $S(B,C)$. Find tuples that agree on the B attribute values and output them.

- **Map:** for each tuple (a, b) from R , output $(b, (\text{R}, a))$ and for each tuple (b, c) from S , produce $(b, (\text{S}, c))$
- **Reduce:** For each input key b , there will a list of values of the form (R, a) or (S, c) . Construct all pairs and output them together with b

Grouping and aggregation


For simplicity, assume we have the relation $R(A,B,C)$ and we want to group-by A and aggregate on B , disregarding C .

- **Map:** for each tuple (a, b, c) from R , output (a, b) . Each key a represents a group.
- **Reduce:** apply the aggregation operator to the list of b values associated with group keyed by a , producing x . Then output (a, x) .

Stage Chaining

- As map reduce calculations get **more complex**, it's useful to break them down into **stages**, with the output of one stage serving as input to the next
- Intermediate output may be useful for **different outputs** too, so you can get some reuse
- The intermediate records can be saved in the data store, forming a **materialized view**
- **Early stages** of map reduce operations often represent the **heaviest amount** of data access, so building and save them once as a basis for many downstream uses saves a lot of work


Matrix Vector Multiplication



The matrix does not fit in memory, and

1. The vector v **does fit** in a machine's memory
2. The vector v **does not fit** in machine's memory


Vector does fit



The matrix is stored in HDFS as a list of (i, j, m_{ij}) tuples


- The elements v_j of v are **available to all mappers**
- **Map:** $((i, j), m_{ij})$ pair $\rightarrow (i, m_{ij}v_j)$ pair
- **Reduce:** $(i, [m_{i1}v_1, m_{i2}v_2, \dots, m_{in}v_n])$ pair $\rightarrow (i, m_{i1}v_1 + m_{i2}v_2 + \dots + m_{in}v_n)$ pair

Vector does not fit



- Divide the vector in equal-sized **subvectors** that can fit in memory
- According to that, divide the matrix in **stripes**
- Stripe i and subvector i are **independent** from other stripes/subvectors
- Use the **previous algorithm** for each stripe/subvector pair

Matrix Matrix Multiplication (I)



Matrix M

Matrix N

Matrix P

$$p_{ik} = \sum_j m_{ij} n_{jk}$$

Matrix Multiplication (II)

- A matrix can be seen as a 3 attributes relation:
 - (row index, column index, value) tuples
 - $M \rightarrow (i, j, m_{ij}), N \rightarrow (j, k, n_{jk})$
- As large matrices are often sparse (0's) we omit such tuples
- The product MN can be seen as a natural join over attribute j , followed by product computation, followed by grouping and aggregation
 - Start with (i, j, v) and (j, k, w)
 - Compute (i, j, k, v, w)
 - Compute $(i, j, k, v \times w)$
 - Compute $(i, k, \sum_j v \times w)$

Matrix Multiplication (III)

- First stage
 - **Map:** given (i, j, m_{ij}) produce $(j, (\text{M}, i, m_{ij}))$
given (j, k, n_{jk}) produce $(j, (\text{N}, k, n_{jk}))$
 - **Reduce:** given $(j, [(\text{M}, i, m_{ij}), (\text{N}, k, n_{jk})])$ produce $((i, k), m_{ij} \times n_{jk})$
otherwise do nothing
- Second stage
 - **Map:** identity
 - **Reduce:** produce the sum of the list of values associated with the key

Matrix Matrix Multiplication (IV)

Algorithm 1: The Map Function

```
1 for each element  $m_{ij}$  of  $M$  do
2   produce (key, value) pairs as  $((i, k), (M, j, m_{ij}))$ , for  $k = 1, 2, 3, \dots$  up
   to the number of columns of  $N$ 
3 for each element  $n_{jk}$  of  $N$  do
4   produce (key, value) pairs as  $((i, k), (N, j, n_{jk}))$ , for  $i = 1, 2, 3, \dots$  up
   to the number of rows of  $M$ 
5 return Set of (key, value) pairs that each key,  $(i, k)$ , has a list with
   values  $(M, j, m_{ij})$  and  $(N, j, n_{jk})$  for all possible values of  $j$ 
```

Algorithm 2: The Reduce Function

```
1 for each key  $(i, k)$  do
2   sort values begin with  $M$  by  $j$  in  $list_M$ 
3   sort values begin with  $N$  by  $j$  in  $list_N$ 
4   multiply  $m_{ij}$  and  $n_{jk}$  for  $j_{th}$  value of each list
5   sum up  $m_{ij} * n_{jk}$ 
6 return  $(i, k), \sum_{j=1} m_{ij} * n_{jk}$ 
```

Matrix Matrix Multiplication (III)

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \times \begin{bmatrix} a & b \\ c & d \\ e & f \end{bmatrix} \longrightarrow \begin{bmatrix} 1a + 2c + 3e & 1b + 2d + 3f \\ 4a + 5c + 6e & 4b + 5d + 6f \end{bmatrix}$$

$(i, k), (M, j, m_{ij})$

$m_{11} = 1$

$(1, 1), (M, 1, 1)k = 1$

$(1, 2), (M, 1, 1)k = 2$

$m_{12} = 2$

$(1, 1), (M, 2, 2)k = 1$

$(1, 2), (M, 2, 2)k = 2$

.....

$m_{23} = 6$

$(2, 1), (M, 3, 6)k = 1$

$(2, 2), (M, 3, 6)k = 2$

$(i, k), (N, j, n_{jk})$

$n_{11} = a \ ((i, k), [(M, j, m_{ij}), (M, j, m_{ij}), \dots, (N, j, n_{jk}), (N, j, n_{jk}), \dots])$

$(1, 1), (N, 1, 1), [(M, 1, 1), (M, 2, 2), (M, 3, 3), (N, 1, a), (N, 2, c), (N, 3, e)]$

$(2, 1), (N, 1, 2), [(M, 1, 1), (M, 2, 2), (M, 3, 3), (N, 1, b), (N, 2, d), (N, 3, f)]$

$(2, 1), [(M, 1, 4), (M, 2, 5), (M, 3, 6), (N, 1, a), (N, 2, c), (N, 3, e)]$

$n_{21} = c \ (2, 2), [(M, 1, 4), (M, 2, 5), (M, 3, 6), (N, 1, b), (N, 2, d), (N, 3, f)]$

$(1, 1), (N, 2, c)i = 1$

$(2, 1), (N, 2, c)i = 2$

$n_{31} = e$

$(1, 1), (N, 3, e)i = 1$

$(2, 1), (N, 3, e)i = 2$

.....

$n_{32} = f$

$(1, 2), (N, 3, f)i = 1$

$(2, 2), (N, 3, f)i = 2$

$[(M, 1, 1), (M, 2, 2), (M, 3, 3), (N, 1, a), (N, 2, c), (N, 3, e)]$

$list_M = [(M, 1, 1), (M, 2, 2), (M, 3, 3)]$

$list_N = [(N, 1, a), (N, 2, c), (N, 3, e)]$

$P(1, 1) = 1a + 2c + 3e$

$P(1, 1) = 1a + 2c + 3e$

$P(1, 2) = 1b + 2d + 3f$

$P(2, 1) = 4a + 5c + 6e$

$P(2, 2) = 4b + 5d + 6f$

Graphs

- $G = (V, E)$, where
 - V represents the set of **vertices** (nodes)
 - E represents the set of **edges** (links)
 - Both vertices and edges may contain **additional information**
- Graph algorithms typically involve:
 - **Performing computations at each node**: based on node features, edge features, and local link structure
 - **Propagating computations**: “traversing” the graph
- Key questions:
 - How do you **represent graph data** in MapReduce?
 - How do you **traverse a graph** in MapReduce?

Representing Graphs (I)

- **Adjacency matrix**
 - Represent a graph as an $n \times n$ square matrix M
 - $n = |V|$
 - $m_{ij} = 1$ means a link from node i to j
- **Advantages:**
 - Amenable to mathematical manipulation
 - Iteration over rows and columns corresponds to computations on outlinks and inlinks
- **Disadvantages:**
 - Lots of zeros for sparse matrices
 - Lots of wasted space


Representing Graphs (II)

- **Adjacency list**
 - Take adjacency matrices...
 - and throw away all the zeros
- Advantages:
 - Much more compact representation
 - Easy to compute over outlinks
- Disadvantages:
 - Much more difficult to compute over inlinks

Shortest Path Algorithm

- Consider simple case of equal edge weights
- Solution to the problem can be defined inductively
- Here's the intuition:
 - Define: b is reachable from a if b is on adjacency list of a
 $\text{DISTANCETO}(s) = 0$
 - For all nodes p reachable from s ,
 $\text{DISTANCETO}(p) = 1$
 - For all nodes n reachable from some other set of nodes M ,
 $\text{DISTANCETO}(n) = 1 + \min(\text{DISTANCETO}(m), m \in M)$

Shortest Path




Shortest Path Algorithm

- Data representation:
 - key: node n
 - value: d (distance from start), adjacency list (list of nodes reachable from n)
 - Initialization: for all nodes except for start node, $d = \text{infinity}$
- Mapper:
 - Selects minimum distance path for each reachable node
 - Additional bookkeeping needed to keep track of actual path
 - adjacency list: emit $(m, d + 1)$
- Sort/Shuffle
 - Groups distances by reachable nodes
- Reducer:
 - Selects minimum distance path for each reachable node
 - Additional bookkeeping needed to keep track of actual path

Details (I)

- Each MapReduce iteration advances the “known frontier” by one hop
 - Subsequent iterations include more and more reachable nodes as frontier expands
 - Multiple iterations are needed to explore entire graph
- Preserving graph structure:
 - Problem: Where did the adjacency list go?
 - Solution: mapper emits (n, adjacency list) as well

Details (II)



Pseudocode

```
1: class MAPPER
2:   method MAP(nid  $n$ , node  $N$ )
3:      $d \leftarrow N.DISTANCE$ 
4:     EMIT(nid  $n$ ,  $N$ )                                ▷ Pass along graph structure
5:     for all nodeid  $m \in N.ADJACENCYLIST$  do
6:       EMIT(nid  $m$ ,  $d + 1$ )                            ▷ Emit distances to reachable nodes

1: class REDUCER
2:   method REDUCE(nid  $m$ , [ $d_1, d_2, \dots$ ])
3:      $d_{min} \leftarrow \infty$ 
4:      $M \leftarrow \emptyset$ 
5:     for all  $d \in \text{counts } [d_1, d_2, \dots]$  do
6:       if ISNODE( $d$ ) then
7:          $M \leftarrow d$                                 ▷ Recover graph structure
8:         else if  $d < d_{min}$  then                      ▷ Look for shorter distance
9:            $d_{min} \leftarrow d$ 
10:     $M.DISTANCE \leftarrow d_{min}$                         ▷ Update shortest distance
11:    EMIT(nid  $m$ , node  $M$ )
```

Graph Algorithm Recipe

- Graph algorithms typically involve:
 - Performing computations at each node: based on node features, edge features, and local link structure
 - Propagating computations: “traversing” the graph
- Generic recipe:
 - Represent graphs as adjacency lists
 - Perform local computations in mapper
 - Pass along partial results via outlinks, keyed by destination node
 - Perform aggregation in reducer on inlinks to a node
 - Iterate until convergence: controlled by external “driver”
 - Don’t forget to pass the graph structure between iterations