

Hadoop

Hadoop Installation

Generally Hadoop can be run in three modes:

1. **Standalone (or local) mode**

- There are no daemons used in this mode.
- Hadoop uses the local file system as a substitute for HDFS file system.
- The jobs will run as if there is 1 mapper and 1 reducer

2. **Pseudo-distributed mode**

- All the daemons run on a single machine
- This setting mimics the behaviour of a cluster
- All the daemons run on your machine locally using the HDFS protocol.
- There can be multiple mappers and reducers (1 reducer by default)

3. **Fully-distributed mode**

- This is how Hadoop runs on a real cluster
- All the daemons run on (a subset of) the cluster's machines using the HDFS protocol
- There are multiple mappers and reducers (1 reducer by default)

Hadoop Program Execution

- Typically, Hadoop programs are developed with a text editors or an IDE
- The actual debugging and execution must be performed on Hadoop without IDE
- Once compiled, an Hadoop program is packaged in a JAR file and submitted to a Hadoop cluster
- On a client, we need to specify where HDFS and Hadoop daemons are running
- For simplicity, we will:
 - Develop programs locally using Maven to manage dependency
 - Package the Hadoop program in a JAR file
 - Copy the file into the Hadoop cluster using SSH
 - Perform debugging and execution on the Hadoop cluster using the Web browser interfaces and the log files
- Be prepared...

Word Count in Hadoop (I)

```
public static class TokenizerMapper extends Mapper<Object, Text, Text, IntWritable>
{
    private final static IntWritable one = new IntWritable(1);
    7 private Text word = new Text();

    4 public void map(Object key, Text value, Context context)
        throws IOException, InterruptedException {
        StringTokenizer itr = new StringTokenizer(value.toString());
        while (itr.hasMoreTokens()) {
            word.set(itr.nextToken());
            context.write(word, one); 6
        }
    }
}
```

Word Count in Hadoop (II)

```
public static class IntSumReducer extends Reducer<Text,IntWritable,Text,IntWritable>
{
    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {

        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}
```

Word Count in Hadoop (I)

```
public static void main(String[] args) throws Exception {  
    Configuration conf = new Configuration(); 1  
    String[] otherArgs = new GenericOptionsParser(conf, args).getRemainingArgs();  
    if (otherArgs.length < 2)  
        System.err.println("Usage: wordcount <in> [<in>...] <out>"); System.exit(2);  
    Job job = Job.getInstance(conf, "word count"); 2  
    job.setJarByClass(WordCount.class); 3  
    job.setMapperClass(TokenizerMapper.class); 4  
    job.setCombinerClass(IntSumReducer.class); 5  
    job.setReducerClass(IntSumReducer.class); 6  
    job.setOutputKeyClass(Text.class); 7  
    job.setOutputValueClass(IntWritable.class); 8  
    for (int i = 0; i < otherArgs.length - 1; ++i)  
        FileInputFormat.addInputPath(job, new Path(otherArgs[i])); 9  
    FileOutputFormat.setOutputPath(job, new Path(otherArgs[otherArgs.length - 1])); 10  
    System.exit(job.waitForCompletion(true) ? 0 : 1); 11  
}
```

Hadoop Terminology

- A MapReduce **job** is a unit of work that the client wants to be performed. It consists of
 - **input** data
 - MapReduce **program**
 - **configuration** information
- Hadoop runs the MapReduce job by dividing it into **tasks**
 - There are two types of tasks: **map tasks** and **reduce tasks**
 - The tasks are **scheduled** using **YARN** (Yet Another Resource Negotiation) and run on nodes in the cluster.
 - If a task fails, it will be automatically rescheduled to run on a different node.
- Hadoop divides the input to a MapReduce job into **fixed-size pieces** called **input splits**
- Hadoop creates **one map task for each split**, which runs the user-defined map function for each **record** in the split

Hadoop Types (I)

Programmers specify **two functions**

- *map function*: from [key (K1), value (V1)] pair to list of [key (K2), value (V2)] pairs
- *reduce function*: from [key (K2), list of values (V2)] pair to list of [key (K3), values (V3)] pairs

```
public class Mapper<K1, V1, K2, V2> {  
    public class Context extends MapContext<K1, V1, K2, V2> {  
        // ...  
    }  
    protected void map(K1 key, V1 value, Context context)  
        throws IOException, InterruptedException {  
        // ...  
    }  
}
```

```
public class Reducer<K2, V2, K3, V3> {  
    public class Context extends ReducerContext<K2, V2, K3, V3> {  
        // ...  
    }  
    protected void reduce(K2 key, Iterable<V2> values, Context context)  
        throws IOException, InterruptedException {  
        // ...  
    }  
}
```


Hadoop Types (II)

- *map function*: from [key (K1), value (V1)] pair to list of [key (K2), value (V2)] pairs
 - *combiner function*: from [key (K2), list of values (V2)] pair to list of [key (K2), values (V2)] pairs
 - *reduce function*: from [key (K2), list of values (V2)] pair to list of [key (K3), values (V3)] pairs
 - *partitioner function*: from [key (K2), value (V2)] pair to integer
- If a combiner function is used, then it has the same form as the reduce function
 - The combiner function is an implementation of Reducer), except its output types are the intermediate key and value types (K2 and V2)
 - Often the combiner and reduce functions are the same, in which case K3 is the same as K2, and V3 is the same as V2
 - The partition function operates on the intermediate key and value types (K2 and V2) and returns the partition index
 - In practice, the partition is determined solely by the key (the value is ignored)

```
public abstract class Partitioner<K2, V2>
{
    public abstract int getPartition(K2 key, V2 value, int numPartitions);
}
```

Hadoop Job Configuration (I)

Property	Job setter method	Input types		Intermediate types		Output types	
		K1	V1	K2	V2	K3	V3
Properties for configuring types:							
mapreduce.job.inputformat.class	setInputFormatClass()	•	•				
mapreduce.map.output.key.class	setMapOutputKeyClass()			•			
mapreduce.map.output.value.class	setMapOutputValueClass()				•		
mapreduce.job.output.key.class	setOutputKeyClass()					•	
mapreduce.job.output.value.class	setOutputValueClass()						•
Properties that must be consistent with the types:							
mapreduce.job.map.class	setMapperClass()	•	•	•	•		
mapreduce.job.combine.class	setCombinerClass()			•	•		
mapreduce.job.partitioner.class	setPartitionerClass()			•	•		
mapreduce.job.output.key.comparator.class	setSortComparatorClass()			•			
mapreduce.job.output.group.comparator.class	setGroupingComparatorClass()			•			
mapreduce.job.reduce.class	setReducerClass()			•	•	•	•
mapreduce.job.outputformat.class	setOutputFormatClass()					•	•

Hadoop Job Configuration (II)

- Java generics' **type erasure** means that the type information isn't always present at runtime, so Hadoop has to be given it explicitly
- Input types are set by the input format.
 - For instance, a `TextInputFormat` generates keys of type `LongWritable` and values of type `Text`
- The other types are set explicitly by calling the methods on the Job
- If not set explicitly, the intermediate types default to the (final) output types, which default to `LongWritable` and `Text`.
 - For instance, if K2 and K3 are the same, you don't need to call `setMapOutputKey Class()`
 - Similarly, if V2 and V3 are the same, you only need to use `setOutputValueClass()`
- Due to type erasure it's possible to configure an Hadoop job with incompatible types, because the configuration isn't checked at compile time.

Hadoop Configuration Defaults

- The only configuration that we set is an **input path** and an output path.
- The **default input format** is `TextInputFormat`
 - The output key is of type `LongWritable`, and the output value is of type `Text`
- The **default mapper** is the `Mapper` class, which writes the input key and value unchanged to the output
 - The output key is of type `LongWritable`, and the output value is of type `Text`
- The **default reducer** is the `Reducer` class, which simply writes all its input to its output
 - The output key is of type `LongWritable`, and the output value is of type `Text`
- The **default partitioner** is `HashPartitioner`, which hashes an intermediate key to determine which partition the key belongs in
 - Each partition is processed by a reduce task
 - So the **number of partitions** is equal to the **number of reduce tasks** for the job
 - **By default there is a single reducer**, and therefore a single partition
- We did not set the **number of map tasks**
 - The number is equal to the **number of splits** that the input is turned into
 - It depends on the size of the input and the file's block size (if the file is in HDFS)

Serialization in Hadoop

- **Serialization** is the process of turning **structured objects** into a **byte stream**
 - for transmission over a network or for writing to persistent storage
- **Deserialization** is the reverse process of turning a **byte stream** back into a series of **structured objects**.
- Serialization is used in two quite distinct areas of distributed data processing
 - for interprocess communication
 - for persistent storage
- In Hadoop, **interprocess communication** between nodes in the system is implemented using **remote procedure calls** (RPCs)
- Hadoop uses its own serialization format called **Writables**
 - compact and fast
 - not so easy to extend or use from languages other than Java
 - There are other serialization frameworks supported in Hadoop, such as **Avro, Thrift, Protobuffers**

Writable Interfaces

```
package org.apache.hadoop.io;
```

```
import java.io.DataOutput;
```

```
import java.io.DataInput;
```

```
import java.io.IOException;
```

```
public interface Writable
```

```
{
```

```
    void write(DataOutput out) throws IOException;
```

```
    void readFields(DataInput in) throws IOException;
```

```
}
```

```
package org.apache.hadoop.io;
```

```
public interface WritableComparable<T> extends Writable, Comparable<T>
```

```
{
```

```
}
```

IntWritable Example (I)

```
package org.apache.hadoop.io;

import java.io.*;

public class IntWritable implements WritableComparable {

    private int value;

    public IntWritable() {}
    public IntWritable(int value) { set(value); }

    public void set(int value) {
        this.value = value;
    }

    public int get() {
        return value;
    }

    public void readFields(DataInput in) throws IOException {
        value = in.readInt();
    }

    public void write(DataOutput out) throws IOException {
        out.writeInt(value);
    }
}
```

IntWritable Example (II)

```
public boolean equals(Object o) {
    if (!(o instanceof IntWritable))
        return false;

    IntWritable other = (IntWritable)o;
    return this.value == other.value;
}

public int hashCode() {
    return value;
}

public int compareTo(Object o) {
    int thisValue = this.value;
    int thatValue = ((IntWritable)o).value;
    return (thisValue < thatValue ? -1 : (thisValue == thatValue ? 0 : 1));
}

public String toString() {
    return Integer.toString(value);
}
```


IntWritable Example (III)

```
/** A Comparator optimized for IntWritable. */
```

```
public static class Comparator extends WritableComparator {
```

```
    public Comparator() {  
        super(IntWritable.class);  
    }
```

```
    public int compare(byte[] b1, int s1, int l1,  
                       byte[] b2, int s2, int l2) {
```

```
        int thisValue = readInt(b1, s1);  
        int thatValue = readInt(b2, s2);
```

```
        return (thisValue < thatValue ? -1 : (thisValue == thatValue ? 0 : 1));
```

```
    }
```

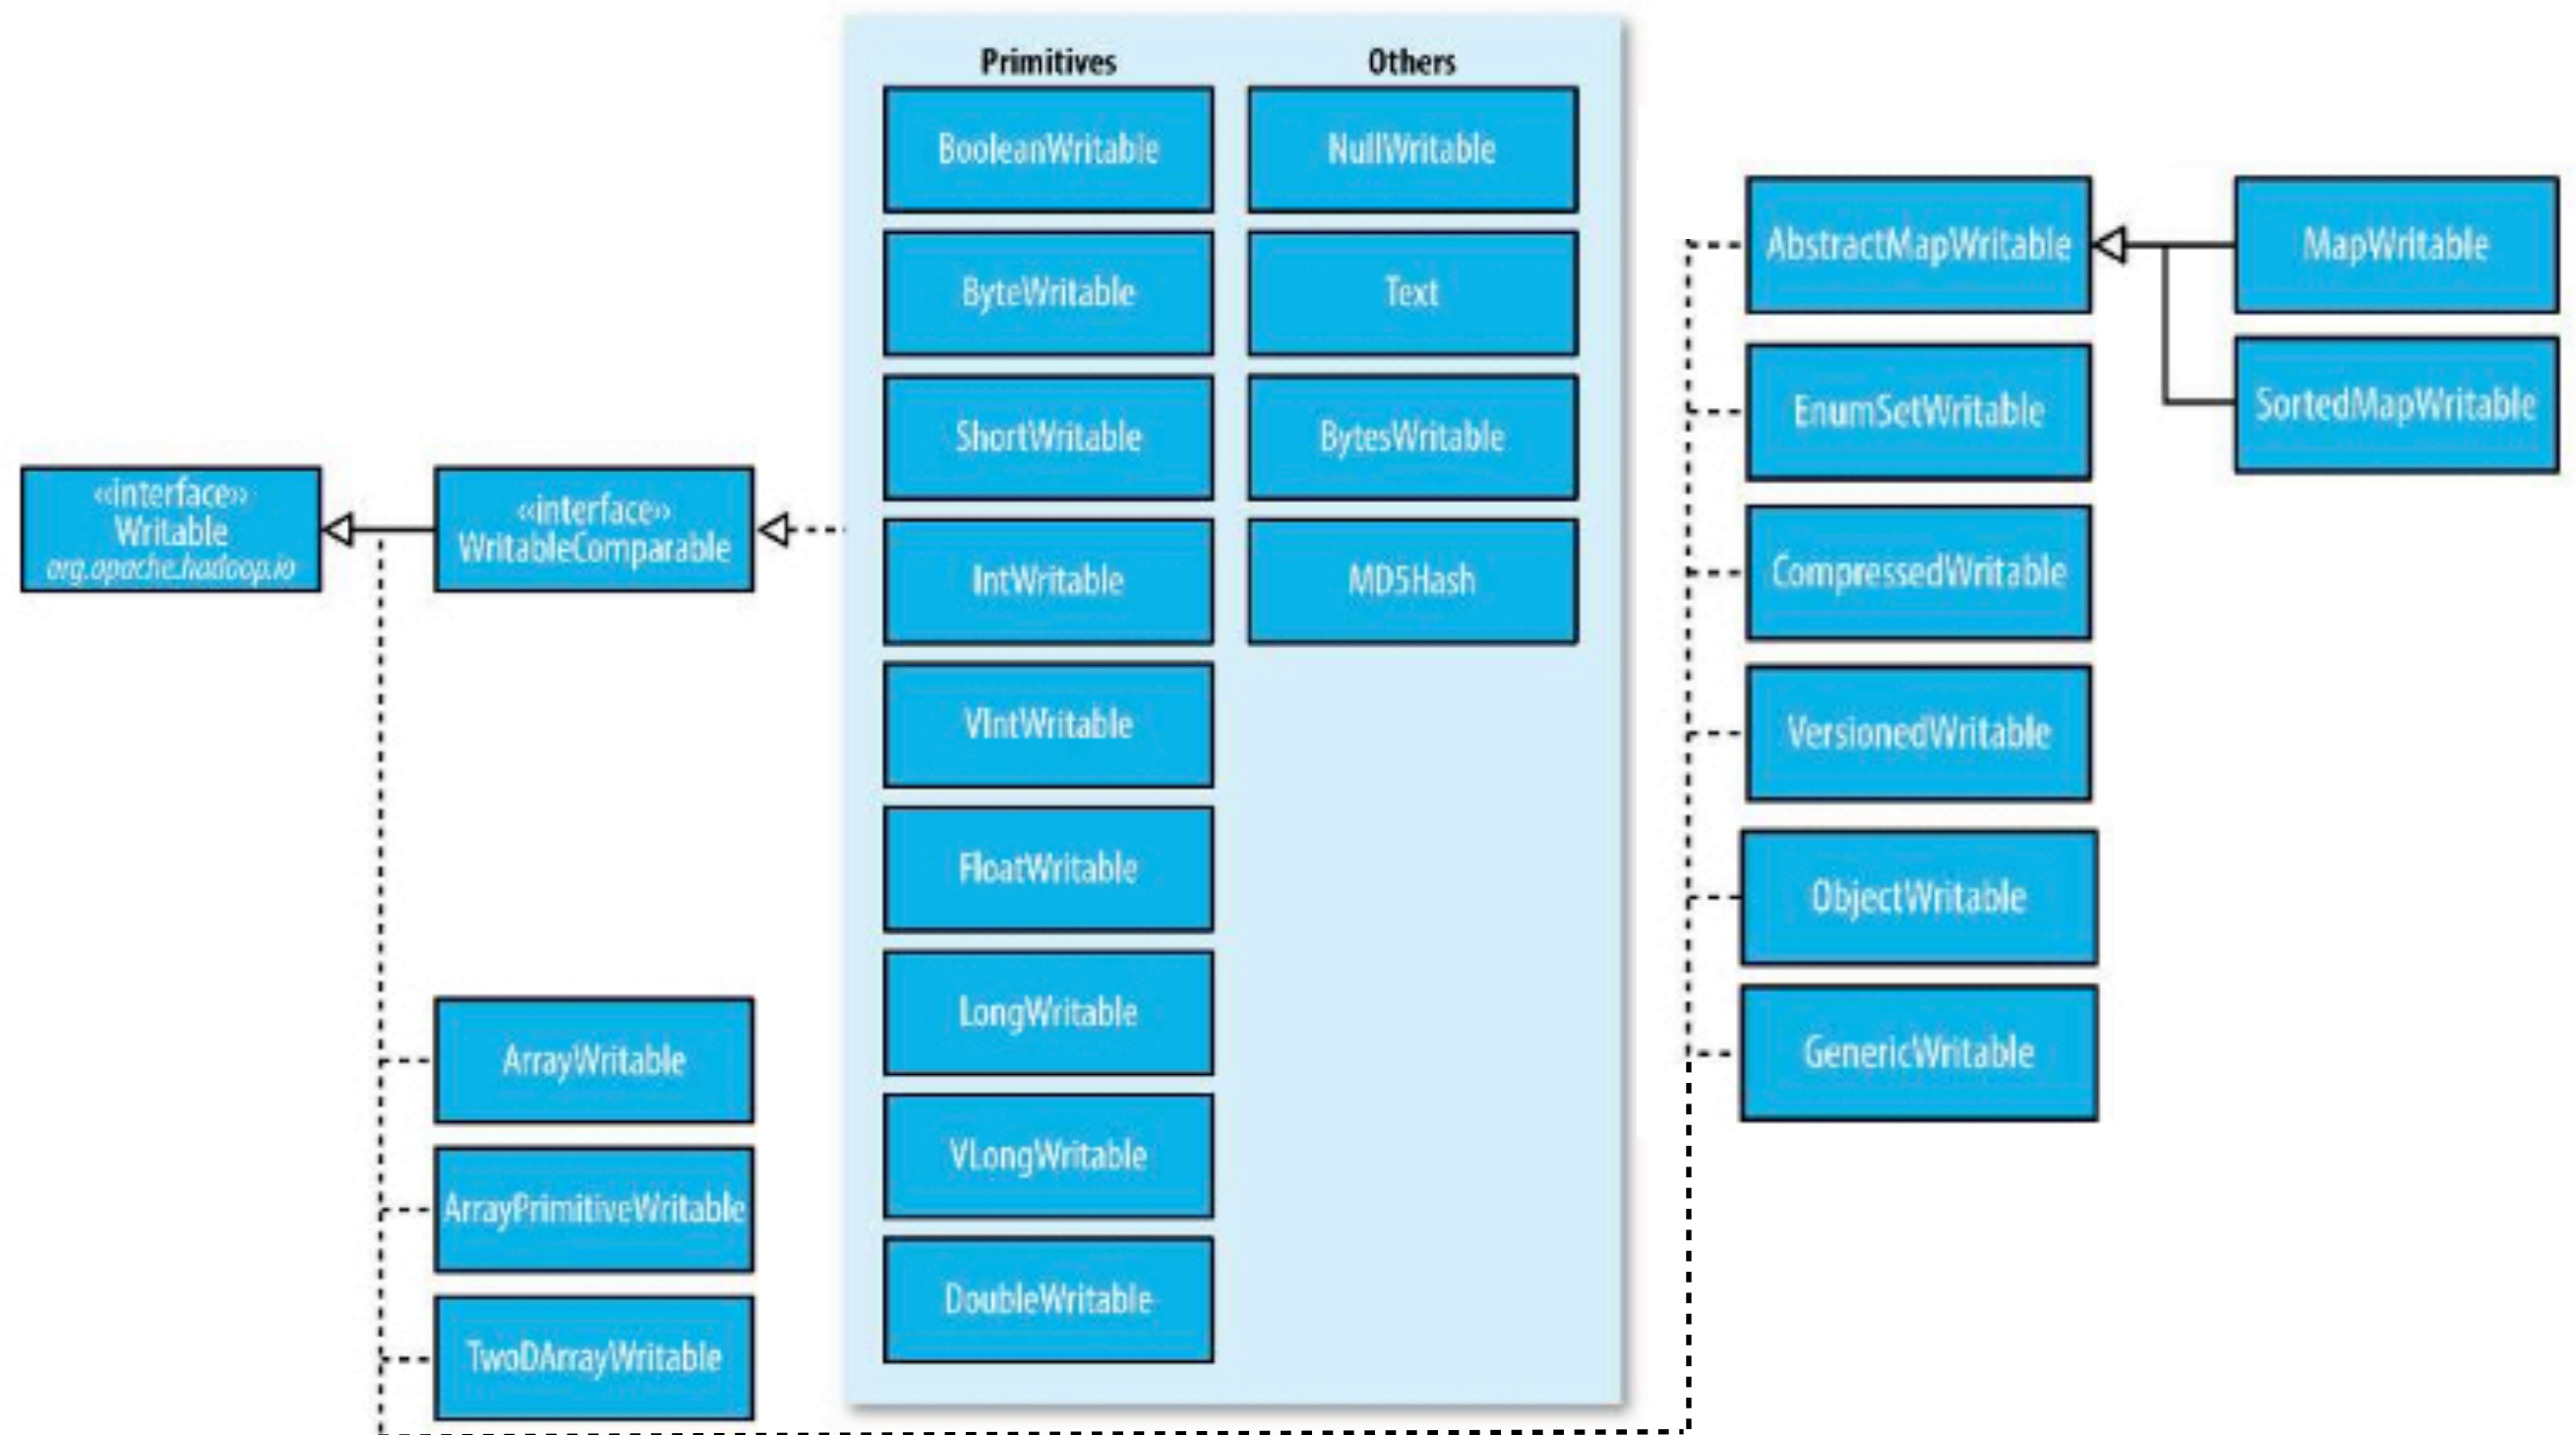
```
}
```

```
}
```

Writable Wrappers

Java primitive	Writable implementation	Serialized size (bytes)
boolean	BooleanWritable	1
byte	ByteWritable	1
short	ShortWritable	2
int	IntWritable	4
	VIntWritable	1–5
float	FloatWritable	4
long	LongWritable	8
	VLongWritable	1–9
double	DoubleWritable	8

Writable Class Hierarchy



Hadoop Input (I)

- An **input split** is a chunk of the input that is processed by a single map task
- Each split is divided into **records**, and the map task processes each record – a key-value pair – in turn
- Splits and records are **logical**
 - Not required to be files, although commonly they are.
 - In a database context, a split might correspond to a range of rows from a table and a record to a row in that range
 - Input splits are represented by the Java class `InputSplit`
- An `InputSplit` has a **length** in bytes and a **set of storage locations** (i.e., hostname strings)
 - A split doesn't contain the input data;
 - A split is just a reference to the data.
 - The **storage locations** are used by Hadoop to **place map tasks** as close to the split's data as possible,
 - The **size** is used to order the splits so that the largest get processed first, to **minimize the job runtime**

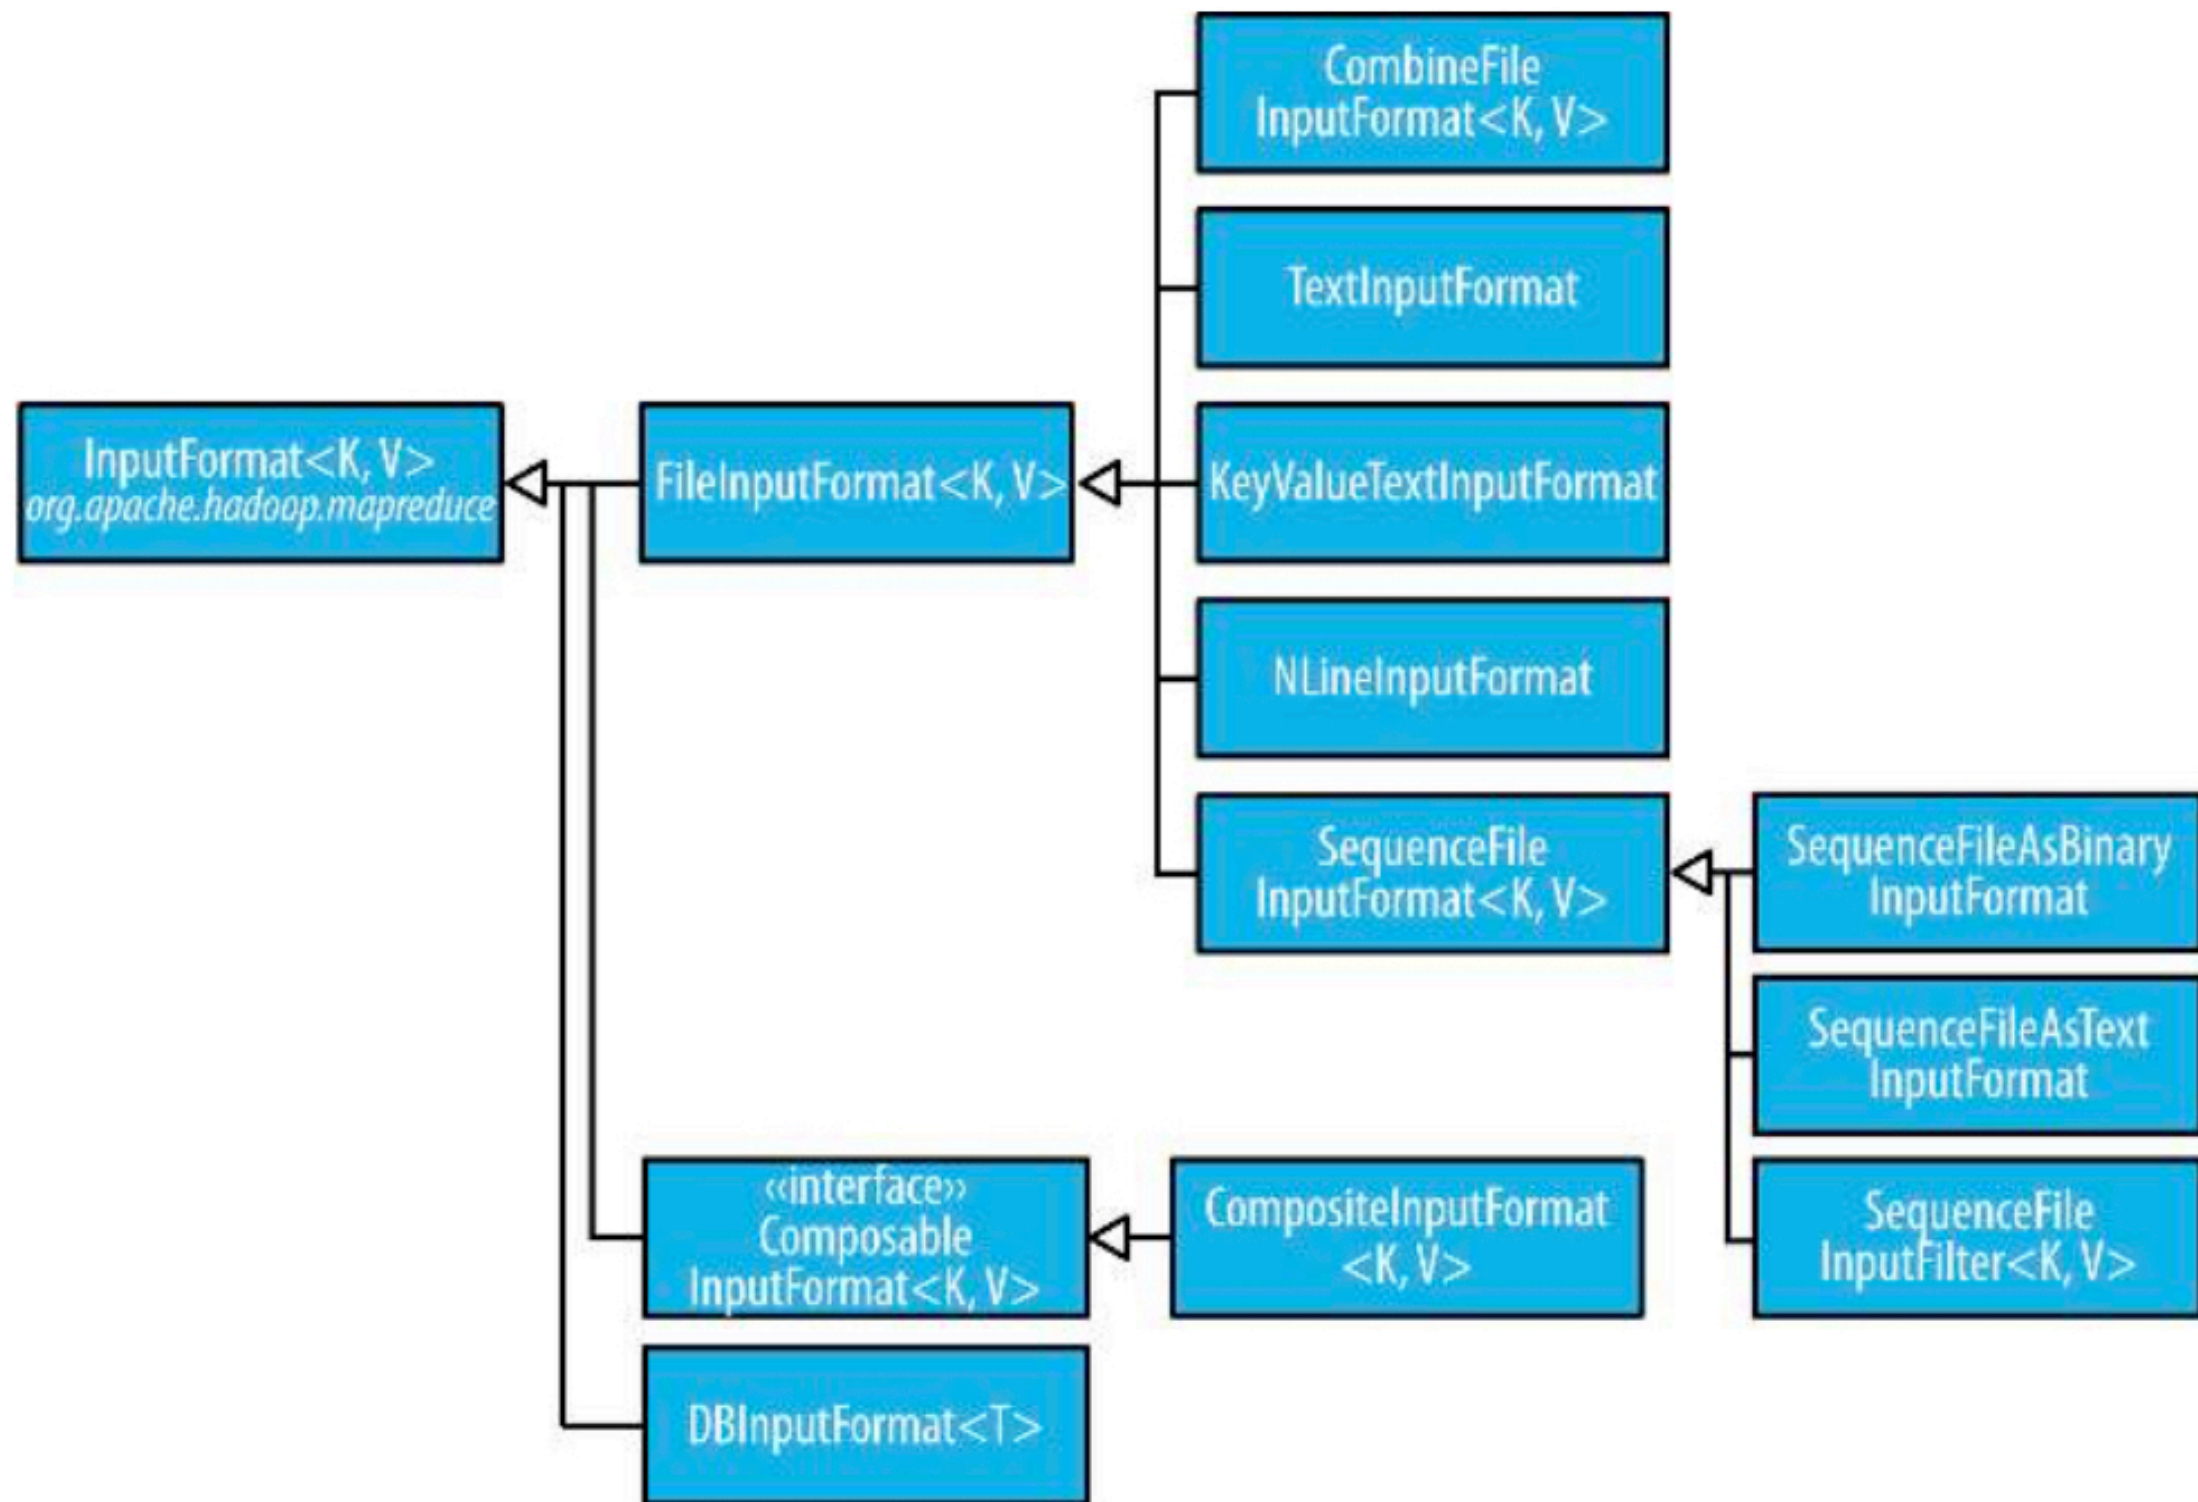
Hadoop Input (II)

- `InputSplits` are created by an `InputFormat` interface implementation

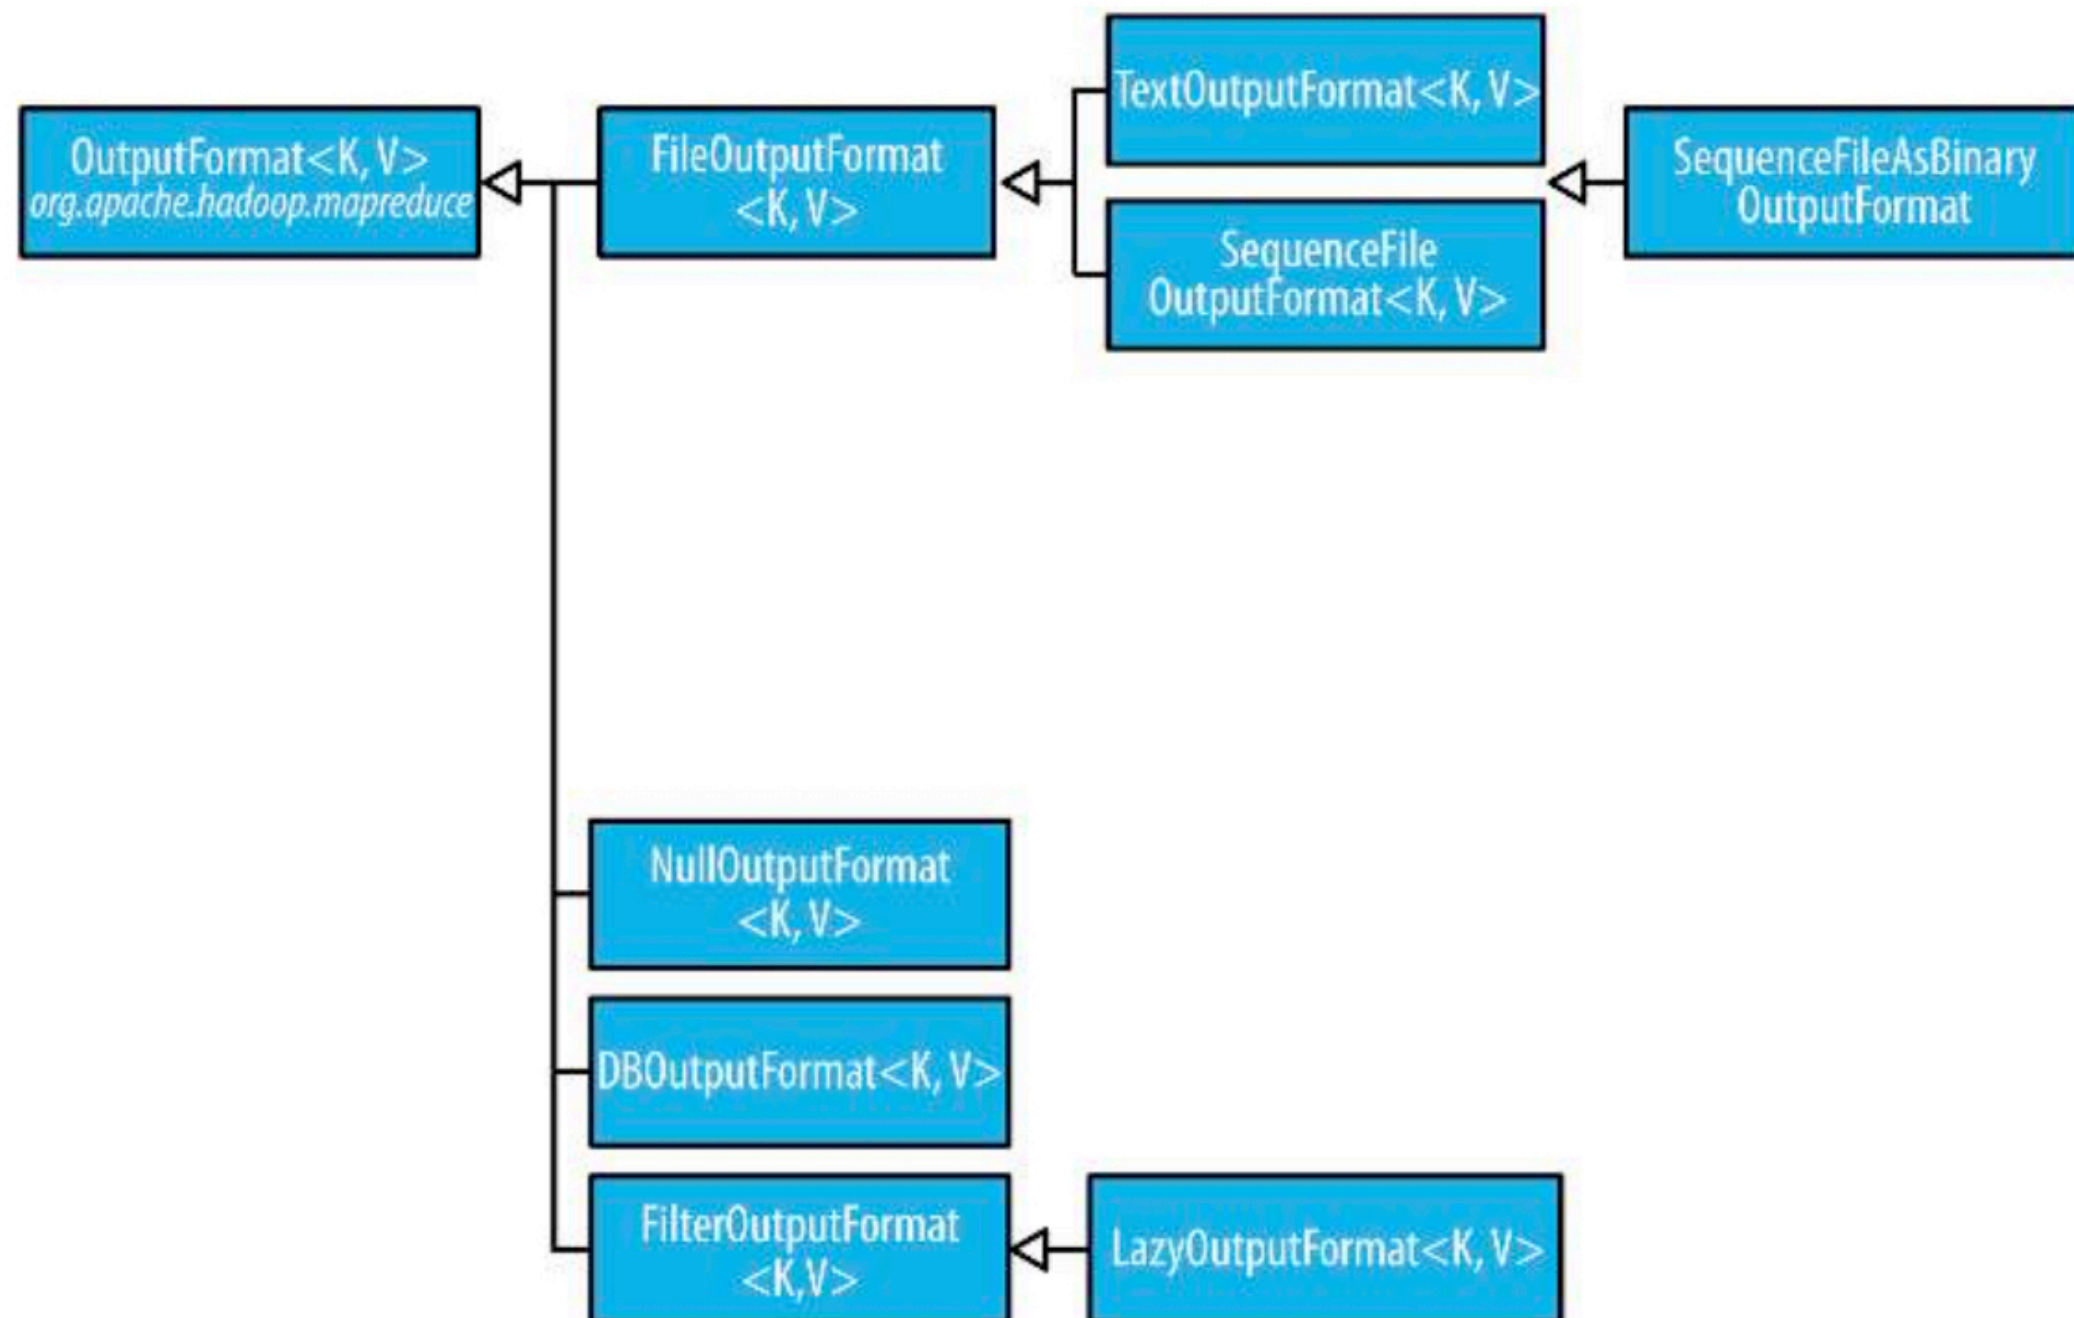
```
public abstract class InputFormat<K, V>
{
    public abstract List<InputSplit> getSplits(JobContext context)
        throws IOException, InterruptedException;
    public abstract RecordReader<K, V> createRecordReader(InputSplit split, TaskAttemptContext context)
        throws IOException, InterruptedException;
}
```

- The client running the job calculates the splits for the job by calling `getSplits()`
- It sends them to the application master, which uses their storage locations to schedule map tasks that will process them on the cluster
- The map task passes the split to the `createRecordReader()` method on `InputFormat` to obtain a `RecordReader` for that split.
- A `RecordReader` is little more than an iterator over records, and the map task uses one to generate record key-value pairs, which it passes to the map function.

InputFormat Class Hierarchy



OutputFormat Class Hierarchy



The setup and cleanup methods

- It is common to want your Mapper or Reducer to execute some code before the `map()` or `reduce()` method is called for the first time
 - Initialize data structures
 - Read data from an external file
 - Set parameters
- The `setup()` method is run before the `map()` or `reduce()` method is called for the first time

```
public void setup(Context context)
```

- Similarly, you may wish to perform some action(s) after all the records have been processed by your Mapper or Reducer
- The `cleanup()` method is called before the Mapper or Reducer terminates

```
public void cleanup(Context context)
```


Passing parameters

```
public class MyDriverClass
{
    public int main(String[] args) throws Exception
    {
        int value = 42;
        Configuration conf = new Configuration();
        conf.setInt ("paramname", value);
        Job job = new Job(conf);
        // ...
        return job.waitForCompletion(true);
    }
}
```

```
public class MyMapper extends Mapper
{
    public void setup(Context context)
    {
        Configuration conf = context.getConfiguration();
        int myParam = conf.getInt("paramname", 0);
        // ...
    }

    public void map...
}
```

Hadoop Tricks

- **Limit as much as possible the memory footprint**
 - Avoid storing reducer values in local lists if possible
 - Use `static final` objects
 - Reuse `Writable` objects
- A **single reducer** is a powerful friend
 - Object fields are **shared** among `reduce()` invocations.
 - The framework reuses value object in reducer, so make **deep copies** if needed
- Passing parameters via class statics doesn't work!
 - Use **configuration parameters** (through Job configuration)
 - Use **external data** sources/sinks (files on HDFS, cache service)
- **An `Iterable<V>` object is not equal to a `List<V>` object**
 - You don't know the number of values **a priori**
 - Accessing an iterable object return a **reference** to the current object
 - Iterable's current object **changes state**, perform a **deep copy** if you need to save it
 - Be careful in allocating a list to store the iterable's elements (memory!)