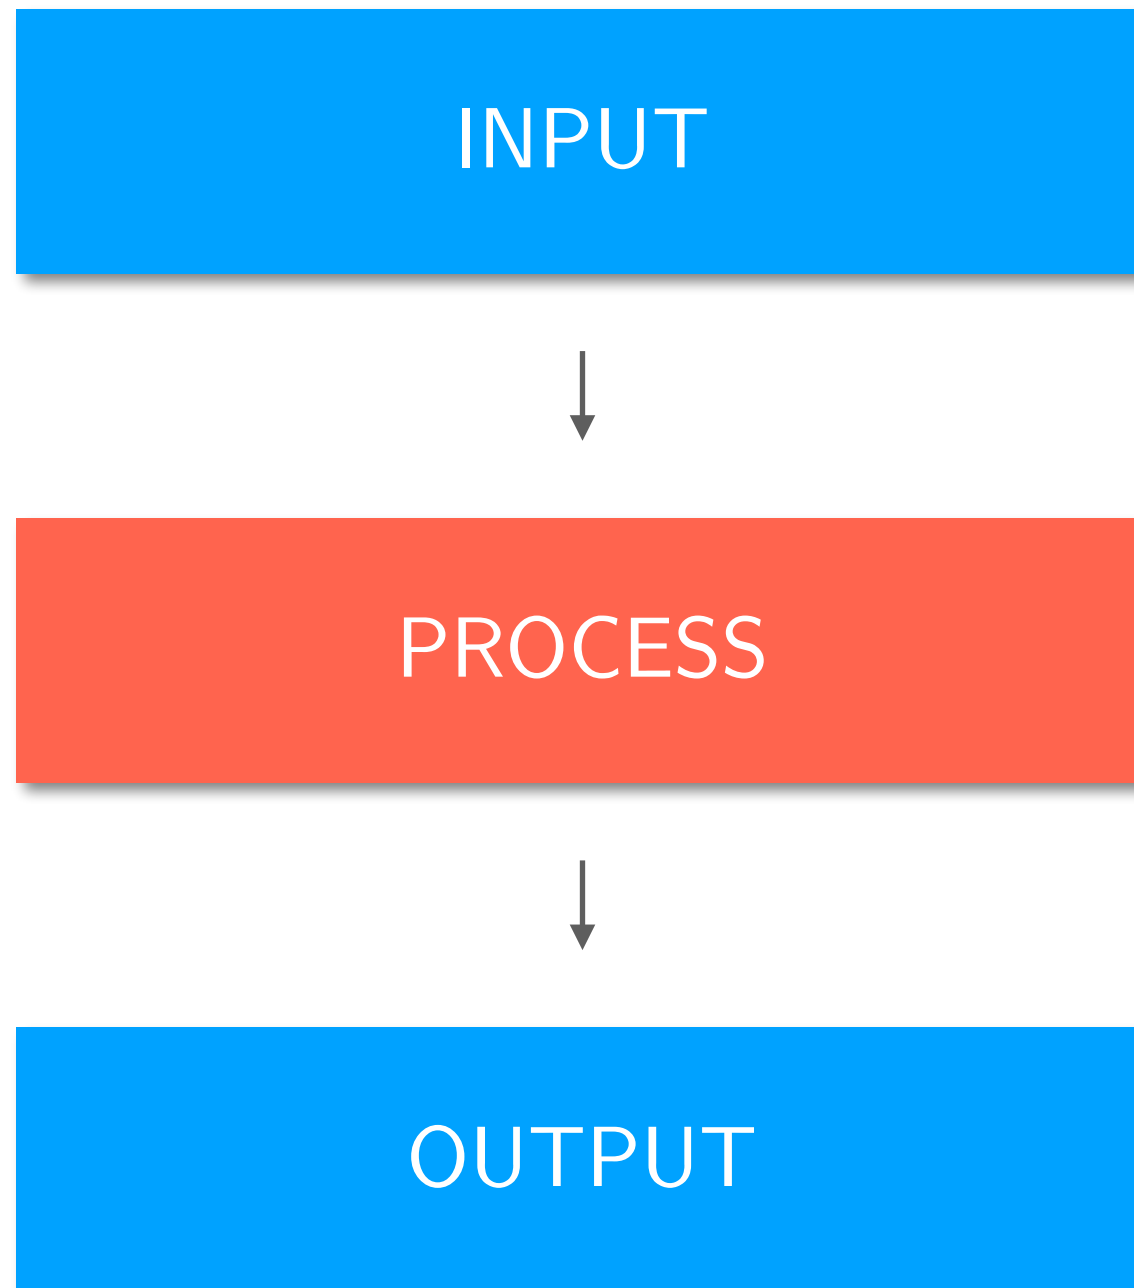


Typical Application



What if?

INPUT




PROCESS



OUTPUT


Parallel Architectures

Shared Memory



- Posix Threads
- OpenMP
- Automatic Parallelization
(Compiler optimizations)

Message Passing



- Sockets
- PVM - Parallel Virtual Machine
(obsolete)
- MPI - Message Passing Interface

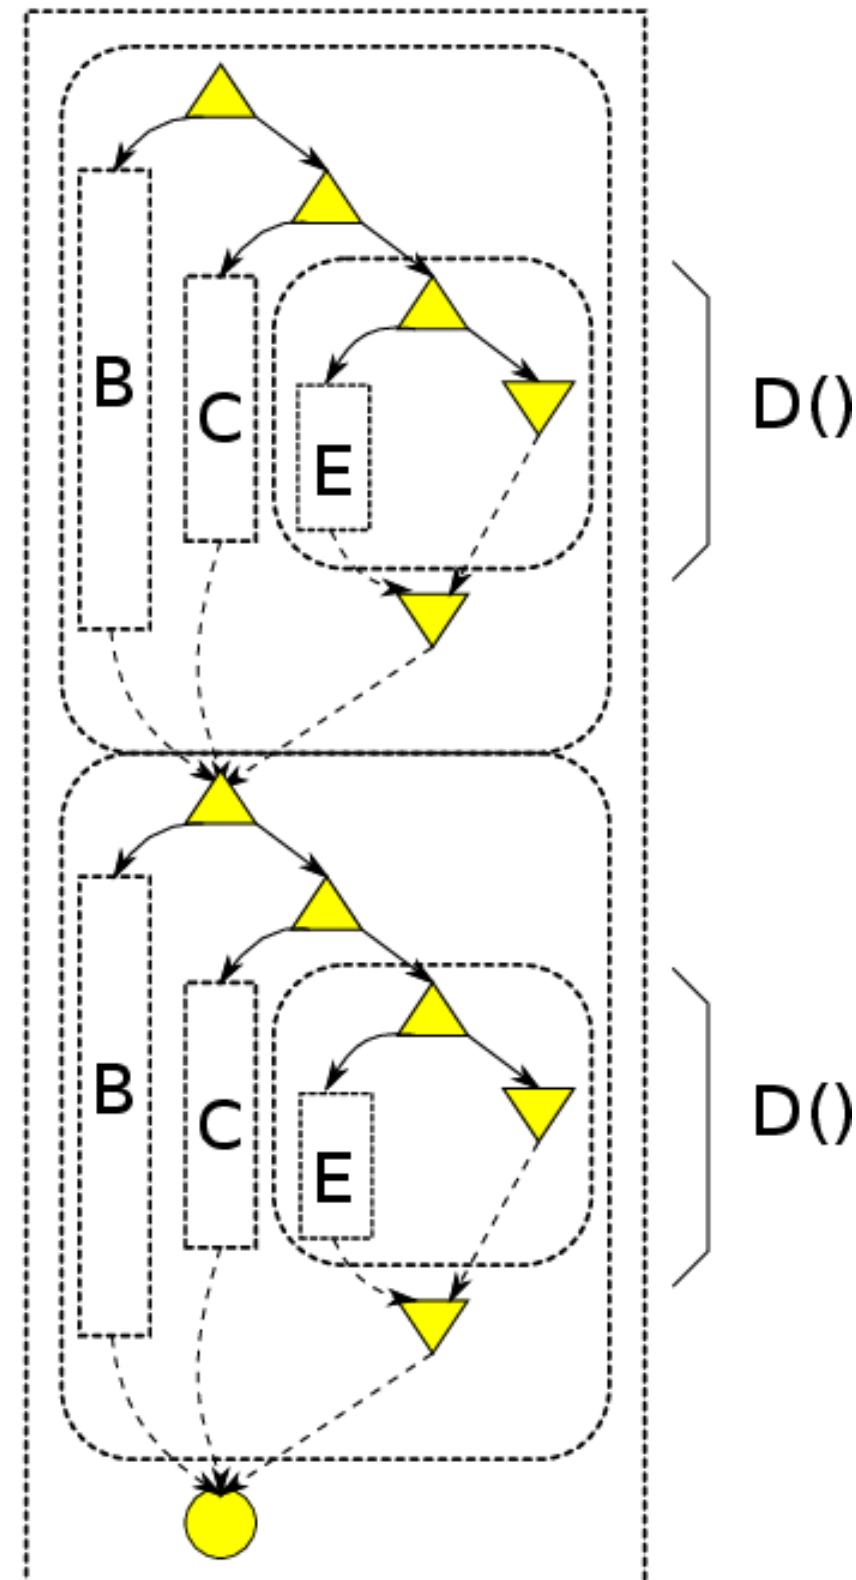

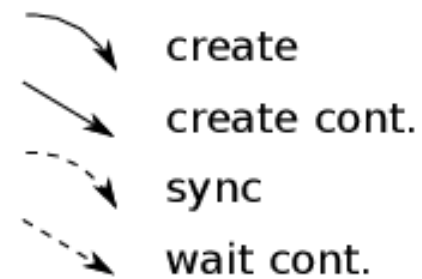

Designing Parallel Algorithms

- Typical steps:
 - Identify what pieces of work can be performed concurrently
 - Partition concurrent work onto independent processors
 - Distribute a program's input, output, and intermediate data
 - Coordinate accesses to shared data: avoid conflicts
 - Ensure proper order of work using synchronization
- Some steps can be omitted
 - For shared memory parallel programming model, there is no need to distributed data
 - For message passing parallel programming model, there is no need to coordinate shared data
 - Processor partition may be done automatically

Task Dependency Graphs

```

A() {
  for (i=0; i<2; i++) {
    CreateTask(B);
    CreateTask(C);
    D();
    WaitTasks();
  }
}
D() {
  CreateTask(E);
  WaitTasks();
}
    
```



Granularity

- Granularity = task size
 - Fine-grain = small tasks, large number of tasks
 - Coarse-grain = large tasks, small number of tasks
 - Choose the proper granularity based on the problem and hardware
- Example: matrix multiplication
 - $N \times N$ matrix A multiply $N \times 1$ vector b give $N \times 1$ vector y
 - Embarrassing parallel: each row of A times vector b gives an element of y
 - Simple decomposition:
 - Task size is uniform
 - No dependencies between task – All tasks share b
 - Fine-grained: each task process one row of A
 - Coarse-grained: each task process three rows of A

Parallelism Level

- Definition
 - number of tasks that can execute in parallel
 - may change during program execution
- Metrics
 - maximum parallelism level: largest # concurrent tasks at any point in the execution
 - average parallelism level: average number of tasks that can be processed in parallel
- Degree of concurrency vs. task granularity
 - inverse relationship

Limits

- What bounds parallel execution time?
 - maximum parallelism degree, e.g. matrix-vector multiplication
example $\leq N^2$ concurrent tasks
 - dependencies between tasks
 - parallelization overheads, e.g., cost of communication between tasks
 - fraction of application work that can't be parallelized
- No single decomposition technique works for all problems

Programmer problems: splitting code

- How to divide code into parallel tasks?
- How to distribute the code?
- How to coordinate the execution?
- How to load the data?
- How to store the data?
- What if more tasks than CPUs?
- What if a CPU crashes?
- What if a CPU is taking too long?
- What if the CPUs are different?
- What if we have a new (serial) code?